

titanic-classification

October 9, 2023

Project Title:

Titanic Survival Prediction System

Problem Statement:

The sinking of the RMS Titanic in 1912 was a tragic event that resulted in a significant loss of life. This project aims to develop a machine learning system that predicts whether a person aboard the Titanic would have survived the disaster based on various factors. The primary goal is to identify the most influential factors that led to survival, such as socio-economic status, age, gender, and more.

Description: The Titanic Survival Prediction System utilizes a dataset containing passenger information, including socio-economic class (Pclass), age, gender, family size, ticket fare, and more.

Conclusion:

Summarize the project's findings and key results. Highlight the factors that were most likely to lead to survival on the Titanic. Discuss the machine learning model's accuracy and its ability to predict survival based on passenger information. Emphasize the importance of understanding historical events like the Titanic disaster through data analysis and predictive modeling.

Libraries

```
[79]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
```

```
[81]: train = pd.read_csv('/content/train.csv')
train.head()
train.tail()
```

```
[81]:
```

	PassengerId	Survived	Pclass	Name \
886	887	0	2	Montvila, Rev. Juozas
887	888	1	1	Graham, Miss. Margaret Edith
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"
889	890	1	1	Behr, Mr. Karl Howell
890	891	0	3	Dooley, Mr. Patrick


	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
886	male	27.0	0	0	211536	13.00	NaN	S
887	female	19.0	0	0	112053	30.00	B42	S
888	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	male	26.0	0	0	111369	30.00	C148	C
890	male	32.0	0	0	370376	7.75	NaN	Q

Exploratory Data Analysis

Missing values

```
[5]: sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

```
[5]: <Axes: >
```




```
[6]: train.isnull().sum().sort_values(ascending=False)
```

```
[6]: Cabin      687
     Age       177
     Embarked    2
     PassengerId 0
     Survived    0
     Pclass     0
     Name       0
     Sex        0
     SibSp      0
     Parch      0
     Ticket     0
     Fare       0
     dtype: int64
```


```
[7]: sns.set_style('whitegrid')
     sns.countplot(x='Survived',data=train,palette='RdBu_r')
```

```
[7]: <Axes: xlabel='Survived', ylabel='count'>
```




```
[8]: sns.set_style('whitegrid')
     sns.countplot(x='Survived',hue='Sex',data=train,palette='RdBu_r')
```

```
[8]: <Axes: xlabel='Survived', ylabel='count'>
```




```
[9]: sns.set_style('whitegrid')  
sns.countplot(x='Survived',hue='Pclass',data=train,palette='rainbow')
```

```
[9]: <Axes: xlabel='Survived', ylabel='count'>
```




```
[10]: train['Age'].hist(bins=30,color='darkred',alpha=0.7)
```

```
[10]: <Axes: >
```




```
[11]: sns.countplot(x='SibSp',data=train)
```

```
[11]: <Axes: xlabel='SibSp', ylabel='count'>
```




```
[12]: sns.countplot(x='Parch',data=train)
```

```
[12]: <Axes: xlabel='Parch', ylabel='count'>
```



```
[13]: train['Fare'].hist(color='green',bins=40,figsize=(8,4))
```


```
[13]: <Axes: >
```



Data Cleaning

```
[14]: plt.figure(figsize=(12, 7))  
      sns.boxplot(x='Pclass',y='Age',data=train,palette='winter')
```

```
[14]: <Axes: xlabel='Pclass', ylabel='Age'>
```




```
[15]: def impute_age(cols):  
      Age = cols[0]  
      Pclass = cols[1]  
  
      if pd.isnull(Age):  
  
          if Pclass == 1:  
              return 37  
  
          elif Pclass == 2:  
              return 29  
  
          else:  
              return 24  
  
      else:  
          return Age
```

```
[16]: train['Age'] = train[['Age', 'Pclass']].apply(impute_age,axis=1)
train['Embarked'] = train['Embarked'].fillna('S')
```

```
[17]: sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

```
[17]: <Axes: >
```



```
[18]: train.drop('Cabin',axis=1,inplace=True)
train.head()
```

```
[18]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

```
Name Sex Age SibSp \
```

0		Braund, Mr. Owen Harris	male	22.0	1
1	Cumings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	1
2		Heikkinen, Miss. Laina	female	26.0	0
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)		female	35.0	1
4		Allen, Mr. William Henry	male	35.0	0

	Parch	Ticket	Fare	Embarked
0	0	A/5 21171	7.2500	S
1	0	PC 17599	71.2833	C
2	0	STON/O2. 3101282	7.9250	S
3	0	113803	53.1000	S
4	0	373450	8.0500	S

```
[19]: train.dropna(inplace=True)
```

Converting Categorical Figures

```
[20]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             891 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Embarked        891 non-null   object
dtypes: float64(2), int64(5), object(4)
memory usage: 76.7+ KB
```

```
[21]: sex = pd.get_dummies(train['Sex'],drop_first=True)
embark = pd.get_dummies(train['Embarked'],drop_first=True)
train.drop(['Sex','Embarked','Name','Ticket'],axis=1,inplace=True)
train = pd.concat([train,sex,embark],axis=1)
train.head()
```

```
[21]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	1	0	3	22.0	1	0	7.2500	1	0	1
1	2	1	1	38.0	1	0	71.2833	0	0	0

2	3	1	3	26.0	0	0	7.9250	0	0	1
3	4	1	1	35.0	1	0	53.1000	0	0	1
4	5	0	3	35.0	0	0	8.0500	1	0	1

Train Test Split

```
[22]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(train.
    ↪drop(['Survived'],axis=1),
                                                    train['Survived'],
    ↪test_size=0.10,
                                                    random_state=101)
```

Training and Predicting

```
[23]: from sklearn.linear_model import LogisticRegression
logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
[23]: LogisticRegression()
```

```
[24]: predictions = logmodel.predict(X_test)
X_test.head()
```

```
[24]:
```

	PassengerId	Pclass	Age	SibSp	Parch	Fare	male	Q	S
331	332	1	45.5	0	0	28.500	1	0	1
700	701	1	18.0	1	0	227.525	0	0	0
748	749	1	19.0	1	0	53.100	1	0	1
751	752	3	6.0	0	1	12.475	1	0	1
481	482	2	29.0	0	0	0.000	1	0	1

```
[26]: predictions
```

```
[26]: array([0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
        1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
        0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0,
        1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
```

$0, 1]$)

```
[82]: X_train_prediction = logmodel.predict(X_train)
      X_train_prediction
```

```
[82]: array([0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0,
0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0,
0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0,
0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0,
1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0,
0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0,
0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,
1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0,
0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0,
0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1,
0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0,
0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0,
0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0,
1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1,
0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1,
1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1,
1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1,
0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1,
0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0,
0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,
1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0,
1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0,
1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 1, 0, 0, 0, 0, 1, 1, 0]])
```

```
[83]: training_data_accuracy = accuracy_score(y_train, X_train_prediction)
print('Accuracy score of training data : ', training_data_accuracy)
```

Accuracy score of training data : 0.8002496878901373

```
[84]: X_test_prediction = logmodel.predict(X_test)
      X_test_prediction
```

```
[84]: array([0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
          1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
          0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0,
          1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
          0, 1])
```

```
[85]: test_data_accuracy = accuracy_score(y_test, X_test_prediction)
      print('Accuracy score of test data : ', test_data_accuracy)
```

Accuracy score of test data : 0.7333333333333333